

25 Years, 2.2M

A data-driven analysis of The New York Times (2000-2025).

Headlines

Revealing shifts in attention, defining eras, and the changing tone of the news.

The Pulse of a Newsroom

The dataset spans a quarter-century of news, covering **2,184,483 usable articles** after dropping corrupted rows and leaked full-text bodies.


Publishing volume peaked in 2006 at over 163,000 articles before the industry shifted toward a "fewer, deeper stories" model.

The newsroom rhythm is remarkably stable: The NYT publishes roughly 6,000 articles daily, dipping only slightly during holidays.

Real analysis requires spotting data artifacts: the sharp 2021 drop reflects an API collection gap, not an editorial cliff.

The Rise and Fall of NYT Article Volume

Annual count of published articles (2000-2025). Peak publication occurred in 2006, while 2022 shows a significant drop.



Source: New York Times Dataset


Words That Define Eras

TF-IDF analysis isolates the unique vocabulary of each news epoch, filtering out structural noise (like obituary boilerplate).

"Iraq" defined the 2000s before vanishing, while "Obama" and "Trump" perfectly mapped to their respective timelines.

"COVID" arrived instantly, dominated the news cycle completely, and faded fast.

Remarkably, **"China"** is the only term consistently present as a steady undercurrent across all 25 years.




Did the News Get Darker?

Automated sentiment analysis reveals a striking trend: the year **2000 was the only period** with a net-positive emotional register.

Every year since has tilted increasingly negative, with the steepest decline starting in 2017 and bottoming out in 2025.

While certain beats (like crime and war) inherently use negative vocabulary, the sustained downward drift over a quarter-century is highly significant.

This highlights a fundamental shift in either editorial tone or the objective reality of the events being covered.



Beyond the Charts

I transformed this 2.2 million row dataset into a fully local
Retrieval-Augmented Generation (RAG) chatbot.

It can answer complex questions about 25 years of coverage using exact
headline citations, with no cloud APIs or data leaks.

If you are a journalist working on stories involving text, media, or messy data at scale,
let's collaborate.

olimiemma.com | linkedin.com/in/olimiemma