

Who dominates the headlines?

An analysis of 2.2 million New York Times records tracks how headline language changes over time.

OLIMI EMMANUEL KASIGAZI

Software and data engineer, New York

Dataset ends December 22, 2025. Coverage is unverified. Rate is approximate.

Independent project. No affiliation with, sponsorship by, or endorsement from The New York Times Company.


≈ 1,500

**titles matching
“Trump”
per 10,000 in 2025**

Highest plotted annual rate
among ten tracked terms.

Trump's 2025 peak stands out

Selected terms at their peak annual rate, per 10,000 filtered article records



A changing mix

Iraq peaks in 2003 and Obama in 2008. Trump reaches the highest plotted rate among ten tracked terms in 2025.


China remains present throughout the period, with a much smaller peak.

Rates compare the share of collected titles matching each term. They do not measure readership or public opinion.

Approximate rates read from supplied trend chart. Case-insensitive substring matching. 2025 is partial.
Source: <https://github.com/olimiemma/NYT-analysis-2K-2K5>

Headline sentiment ends at its lowest point

Average VADER lexical sentiment by year, across the filtered corpus



-0.0494

2025 annual mean

The overall trend falls, with several recoveries. Both 2000 and 2009 have positive means.


A lexical measure

VADER scores word associations. Topic mix and missing records can affect the trend. This does not establish editorial bias or a causal effect of political events.

Annual means only, from saved notebook output. 2,071,261 filtered records. 2025 is partial.
Source: <https://github.com/olimiemma/NYT-analysis-2K-2K5>

The most frequent word was “paid”

Repeated publication language can overwhelm a simple headline ranking



“Paid” and “notice” reflect repeated notice language.
Selected counts exclude common stopwords.
Recurring templates can obscure topical patterns.

Word occurrences, not unique articles. Annual volume uses an earlier 2,216,466-record stage. Coverage is unverified.
Source: <https://github.com/olimiemma/NYT-analysis-2K-2K5>

What the counts include

2,184,483

cleaned records for word counts

2,071,261

after removing 113,222 records with boilerplate titles

The separate entity analysis processed 399,988 titles, against a 400,000 target. Ambiguous names can merge people or brands.

Three stories worth reporting

An independent project with inspectable code and questions for further reporting

01 Who dominates the news agenda?

Trump's peak stands out among tracked terms. How does that pattern compare across outlets and news sections?

02 How much can sentiment scores tell us?

The annual mean declines overall. A manual review could test how much reflects topic mix or scoring errors.

03 What disappears when we clean an archive?

Boilerplate changes the analytical corpus by 113,222 records. Filtering choices deserve scrutiny alongside the findings.

Contact the author

**OLIMI EMMANUEL
KASIGAZI**

Software and data engineer
New York

olimiemma@gmail.com
<https://olimiemma.com/>

Available to discuss the findings,
methodology and potential
coverage.

Code and methodology: <https://github.com/olimiemma/NYT-analysis-2K-2K5>

Experimental retrieval chatbot: designed for local models. The repository does not establish a completed offline run.